

Does Narrow Deontological IPD Supervision Transfer? A Norm-vs-Heuristic Test

Dan Le

Mini-project for CAISH Mars V / Agentic Moral Alignment

Question

Tennant's *Moral Alignment for LLM Agents* studies whether intrinsic moral rewards can shape the actions of LLM agents in simple game environments. Her CAISH project description asks whether policies learned in these settings generalize beyond the training domain and whether they are robust to persona-prompting. I test two narrow questions. First, if Gemma2-2b-it is supervised on the deontologically preferred action in IPD states where the opponent previously cooperated, does it learn a conditional norm, overgeneralize into unconditional cooperation, or merely learn a surface action-token policy? Second, whatever policy this SFT installs, is it more robust to adversarial persona prompts than the base model?

This is not a PPO replication. I use QLoRA SFT as a rapid baseline for the evaluation question in Tennant's Project Suggestion #3: before setting up a full RL pipeline, can a compact behavioral-cloning approximation to the deontological action signal install any policy that is robust to adversarial framing? The training target approximates the action-level signal that Tennant's deontological PPO reward would prefer in cooperative-prior states.

Method

I fine-tuned Gemma2-2b-it on abstract Iterated Prisoner's Dilemma prompts using arbitrary action tokens. Training examples included only states where the previous opponent action was the cooperative token, action1. In those states, Tennant's deontological norm gives a clear prescription: do not defect against a prior cooperator. The target completion was always action1.


I call this **deon-preference SFT** to emphasize that it supervises only the action label the deontological reward would prefer in cooperative-prior states; it is not the full intrinsic reward signal or PPO training process. This setup intentionally creates a possible failure mode: the model may learn "always output action1" rather than a conditional norm. I test that directly by evaluating held-out prior-defection states.

I compare the base and tuned models on three evaluation suites: original action1/action2 IPD prompts, new-token action3/action4 IPD prompts, and 20 natural-language reciprocal dilemmas with cooperative labels balanced between A and B. To test conditionality, I compare cooperation after prior cooperation versus prior defection.

I also run a small persona-prompting stress test on the natural-language dilemmas using neutral, ruthless-game-theorist, tournament-maximizer, villain-roleplay, enemy-framing, and authority-pressure prompts.

Results

The clearest positive result came from the persona stress test. The tuned model showed smaller cooperation drops from its own neutral baseline under ruthless-game-theorist prompting (tuned -0.20 vs. base -0.55), villain-roleplay (tuned -0.30 vs. base -0.45), and authority-pressure (tuned -0.05 vs. base -0.15). Enemy/outgroup framing was the exception: both models collapsed, with base cooperation at 0.05 and tuned cooperation at 0.10.



Persona average cooperation

This robustness did not come from learning the intended conditional deontological norm. In the original action1/action2 IPD suite, the tuned model chose the cooperative token after both prior cooperation and prior defection: $P(C \text{ given prior coop})=1.00$, $P(C \text{ given prior defect})=1.00$. I interpret this as unconditional cooperation, not conditional norm learning.

Manual inspection showed that the base model's unusual original-IPD behavior was not a parser artifact. Base Gemma2-2b-it produced clean action1 / action2 outputs, but its choices were highly sensitive to minor prompt variants and previous self-action. For example, when $\text{prev_self}=\text{action2}$ and $\text{prev_opp}=\text{action1}$, it chose action2 in all five prompt variants. I therefore treat base original-IPD behavior as a prompt-sensitivity diagnostic rather than a stable reciprocal-policy estimate.

The remaining transfer suites mostly exposed ceiling effects. New-token IPD was uninformative because both base and tuned models chose the cooperative token 100% of the time. Neutral natural-language dilemmas were also near ceiling for the base model, so they did not provide clean evidence of conditional transfer.

Interpretation

The main negative result is clear: narrow deon-preference SFT did not recover Tennant's conditional deontological rule. It produced a broad cooperation heuristic. The more interesting positive result is that this heuristic was more resistant than base behavior to several adversarial persona prompts, measured as cooperation drop from each model's own neutral baseline.

This is preliminary evidence, not a definitive robustness result. Each persona condition contains only 20 hand-written scenarios, and I trained only one deon-preference SFT model. Without selfish-target, utilitarian-target, or generic cooperative SFT controls, I cannot tell whether the robustness gain is specific to the deontological target or would arise from any training procedure that repeatedly reinforces cooperative actions.

Next steps

1. Add selfish-target, utilitarian-target, and generic cooperative SFT controls to test whether persona robustness is specific to the moral target.
2. Repeat the same evaluation on Tennant-style PPO checkpoints to separate effects of the training algorithm from effects of the target behavior.
3. Build harder natural-language dilemmas that avoid base-model ceiling effects and better distinguish conditional reciprocity from generic prosociality.

Items 1 and 2 are the natural first month of work if accepted: the control SFT targets are cheap to train, and the persona evaluation suite already exists for comparison against Tennant-style PPO checkpoints.